

Machine Learning Engineer

Location: Mumbai (Andheri East), India (In-Office)

Experience: 3+ Years

About Skima Innovation:

At Skima, we don't just build models; we build the future. We are a dynamic team dedicated to pushing the boundaries of what's possible through data-driven innovation. We are looking for a talented Machine Learning Engineer who is ready to take ownership of end-to-end ML lifecycles and transform complex data into scalable, real-world solutions.

The Role:

As an ML Engineer at Skima, you will sit at the intersection of data science and software engineering. You won't just be "playing with data"—you will be designing, developing, and deploying high-performance models that drive our core products. You will work in a collaborative environment where your algorithms directly impact business outcomes.

Key Responsibilities

- **Production Pipelines:** Architect and manage automated ML implementation pipelines for seamless transition from research to production.
- **Deep Learning Deployment:** Optimize and deploy large-scale Deep Learning models using specialized inference engines.
- **Containerization & Orchestration:** Package ML services using Docker and manage deployments via Kubernetes to ensure high availability and scalability.
- **MLOps Mastery:** Establish CI/CD for ML, implementing automated testing, versioning (DVC), and model registry workflows.
- **Model Observability:** Implement comprehensive monitoring for model drift, data integrity, and real-time performance latency.
- **Optimization:** Fine-tune models for resource efficiency, focusing on quantization and pruning for production-grade inference.

What You Bring

- **Experience:** 3+ years of hands-on experience in ML engineering with a focus on production-grade deployments.
- **MLOps Stack:** Proficiency with tools like MLflow, Kubeflow, W&B for managing the model lifecycle.
- **Cloud & Infrastructure:** Strong experience with AWS/Azure/GCP ML services and containerized environments.
- **Technical Depth:** Expert-level Python and deep familiarity with PyTorch or TensorFlow.

- Engineering Rigor: Solid understanding of software design patterns, API development (FastAPI/Flask), and database management.
- Soft Skills: Ability to translate complex technical roadmaps into clear business value for stakeholders.
- Familiarity with tuning vLLM based inference is a plus.

Education & Qualifications

Bachelor's or Master's degree in Computer Science, Data Science, Mathematics, or a related technical field.